


SegNet for Semantic Segmentation

- SegNet is a deep learning network with encoder-decoder architecture for semantic segmentation and was researched and developed by members of the Computer Vision and Robotics Group at the University of Cambridge.
- SegNet focuses on capturing and storing boundary information from encoder feature maps before sub-sampling. The more efficient way to store information is to store only the max-pooling indices, i.e, the locations of the maximum feature value in from feature map with a pooling window.

More information in Summary of literature survey [here](#)

****Implementation Detail of Segnet:****



- SegNet has an encoder network and a corresponding decoder network, followed by a final pixel-wise classification layer.
- The encoder network consists of 13 convolutional layers that correspond to the first 13 convolutional layers in the VGG16 network designed for object classification and hence can initialize pre-trained weights.
- Each encoder in the encoder network performs convolution with a filter bank to produce a set of feature maps that are batch normalized with ReLU.
- Followed by Max-pooling with a 2x2 window and stride 2 to get sub-sampled output by a factor of 2 where meanwhile max pool indices are memorized that is to capture and store boundary information of the feature maps.
- The Decoder network upsamples its input feature map using the memorized max-pooling indices from the corresponding encoder feature map which results in sparse feature map and less number of parameters when compared to other architectures.
- Followed by feature maps are then convolved with a trainable decoder filter bank to produce dense feature maps with batch normalization.
- The final decoder output is fed to a multi-class softmax classifier to produce class probabilities for each pixel independently.
- For more information please refer the paper [SegNet Deep Convolutional Encoder-Decoder Architecture for Image Segmentation](#)

Importing required Modules

In []:

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import backend as K
from numpy import asarray, zeros, moveaxis
from tensorflow.keras.initializers import *
from tensorflow.keras.models import *
from tensorflow.keras.layers import *
from tensorflow.keras.callbacks import *
from tensorflow.keras.optimizers import *
import matplotlib.pyplot as plt
from sys import getsizeof
from keras.callbacks import TensorBoard
from tqdm import tqdm_notebook, tqdm
from sklearn.metrics import *
import os, sys, ntpath, fnmatch, shutil, cv2
import joblib, os.path, itertools, warnings
from keras.models import load_model
from scipy.sparse import csc_matrix
import numpy as np
import pandas as pd
from os import path
!pip install import_ipynb
from IPython.display import clear_output
from time import time
np.random.seed(0)
from google.colab import drive
drive.mount('/content/drive')
```

```

❗ cp -v "/content/drive/My Drive/IID_Files/Utility_Functions.ipynb" "/content"
❗ cp -v "/content/drive/My Drive/IID_Files/Final.ipynb" "/content"
❗ cp -v "/content/drive/My Drive/IID_Files/IID_Data_Prep_Utils.ipynb" "/content"
warnings.filterwarnings("ignore")
clear_output()

```

Importing Data Preparation Modules

In []:

```

import import_ipynb
from Utility_Functions import *
from Final import *
from IID_Data_Prep_Utils import *

```

Checking Status:

```

-----
1.Image Data Preparation      .. .. .. >>> |Done| <1/5>
2.Label Mask Preparation     .. .. .. >>> |Done| <2/5>
3.Data Shuffling             .. .. .. >>> |Done| <3/5>
4.Data Train_Test_Split      .. .. .. >>> |Done| <4/5>
5.Loading Final Data         .. .. .. >>> |Done| <5/5>
-----

```

Gen RAM Free: 2.41 GB - Used: 23.22 GB - Total : 25.51 GB - Util 91.02 %
GPU RAM Free: 15.9 GB - Used: 0.0 GB - Total : 15.9 GB - Util 0.0 %

General Utility Function for Prediction

In []:

```

def predict_for(data_for_prediction, weights_save_path=False):

    """ General Function to perform prediction for the specified data split """

    Mean_MIoU, Accuracy, cf_matrix=[], [], np.zeros((7,7))
    x, y = Load_For_Prediction(data_for_prediction)
    Model, Skip = Select_Model(weights_save_path), 2

    for d in tqdm_notebook(range(0, len(x), Skip)):

        if (d>=(len(x)-Skip)):
            plot,_,_=True,clear_output(),print("Total number of samples in {0} : {1}".format(data_for_prediction, len(x)))
            else: plot=False

            Miou, cf_matrix, Accuracy=Function_2(x[d:d+Skip],y[d:d+Skip],Mean_MIoU,cf_matrix,Accuracy,Model,
            plot,False,False,False)
            collected = gc.collect()

    return Miou, Accuracy, cf_matrix

```

Segnet Maxpooling and Unpooling with indices

In []:

```

from keras.layers import *
from keras.models import *
import numpy as np

def MaxUnpooling2D(pool, ind, batch_size, name):

    """
    Function to perform UpSampling with Indices
    Input  : pool <tf.keras.layer>, ind <tf.keras.layer>, batch_size <Int>, name<String>
    Return : ret <tf.keras.layer> """

    # Reference:https://stackoverflow.com/questions/60104102/custom-max-pool-layer-valueerror-the-channel-dimension-of-the-inputs-should-be

    with tf.compat.v1.variable_scope(name):

        # Computing Upsampling Output Shape
        output_shape=[None,pool.shape[1]*2,pool.shape[2]*2,pool.shape[3]]

        # Flattening the pool tensor
        pool_ = tf.reshape(pool, [-1])

```

```

# creating tensors to arrange indicies in batch wise
batch_range = tf.reshape(tf.range(batch_size,dtype=ind.dtype), [tf.shape(pool)[0], 1, 1, 1])
b = tf.reshape(tf.ones_like(ind) * batch_range, [-1, 1])

# Flattening the index tensor concating with tensor
ind_ = tf.concat([b, tf.reshape(ind, [-1, 1])], 1)

# Preforming Upsampling on pool tensor values with indicies
ret = tf.scatter_nd(ind_, pool_, shape=[batch_size, output_shape[1] * output_shape[2] * output_shape[3]])

# Reshaping tensor according to Sample Size
ret = tf.reshape(ret, [tf.shape(pool)[0], output_shape[1], output_shape[2], output_shape[3]])

return ret

```

Implementation of Segnet

In []:

```

def Segnet_Segmentation(input_shape, n_labels, batch_size, Kernel=(3,3), Pool_Filter=(2,2), output_mode=
"softmax"):

    """
    Function to build Segnet Architecture for Image Segmentation
    Input : input_shape <Tuple>, n_labels <Int>, batch_size <Int>, Kernel <Tuple>, Pool_Filter <Tuple>, o
utput_mode <String>
    Return : model """

    # Defining input_shape and Batch Size
    Inputs, batch_size = Input(shape=input_shape), batch_size

def Segnet_Encoder(Block_Number, Filters, Input_layer):

    """
    Function to Build Segnet Encoder Blocks
    Input : Block_Number <Int>, Filters <Int>, Input_layer <Keras.layer>
    Return : Layer_2 <Keras.layer> """

    # Defining First Convolution layers with Batch Normalization for each Encoder Block
    Layer_1 = Convolution2D(Filters, Kernel, name= "En_Block"+str(Block_Number)+"_Conv1", activation
= 'relu', padding= "same", kernel_initializer= "he_normal")(Input_layer)
    Layer_1 = BatchNormalization(name= "En_Block"+str(Block_Number)+"_Batch1")(Layer_1)

    # Defining Second Convolution layers with Batch Normalization for each Encoder Block
    Layer_2 = Convolution2D(Filters, Kernel, name= "En_Block"+str(Block_Number)+"_Conv2", activation
= 'relu', padding= "same", kernel_initializer= "he_normal")(Layer_1)
    Layer_2 = BatchNormalization(name= "En_Block"+str(Block_Number)+"_Batch2")(Layer_2)

    if Block_Number>2:
        # Defining Third Convolution layers with Batch Normalization for necessary Encoder Block

        Layer_3 = Convolution2D(Filters, Kernel, name= "En_Block"+str(Block_Number)+"_Conv3", activa
tion = 'relu', padding= "same", kernel_initializer= "he_normal")(Layer_2)
        Layer_3 = BatchNormalization(name= "En_Block"+str(Block_Number)+"_Batch3")(Layer_3)

        if Block_Number==5:
            # Zero Padding of Feature map
            Layer_3=ZeroPadding2D((1,0),(0,0)), name='Zero_pad')(Layer_3)
            return Layer_3

    return Layer_2

def Segnet_Decoder(Block_Number, Filters, Input_layer):

    """
    Function to Build Segnet Decoder Blocks
    Input : Block_Number <Int>, Filters <Int>, Input_layer <Keras.layer>
    Return : Layer_2 <Keras.layer> """

    # Defining Filters based on Decoder Block
    Get_filter = lambda x : int(Filters/2) if ((Block_Number > 2 and Block_Number < 5) and (x==3))
or (Block_Number==2 and x==2)) else Filters

    if Block_Number==5:
        # Cropping of Feature map
        Input_layer=Cropping2D((1, 0),(0,0))(Input_layer)

    # Defining First Convolution layers with Batch Normalization for each Decoder Block

```

```

Layer_1 = Convolution2D(Get_filter(1), Kernel, name= "Dec_Block"+str(Block_Number)+"_Conv1", activation = 'relu', padding= "same", kernel_initializer= "he_normal")(Input_layer)
Layer_1 = BatchNormalization(name= "Dec_Block"+str(Block_Number)+"_Batch1")(Layer_1)

# Defining Second Convolution layers with Batch Normalization for each Decoder Block
Layer_2 = Convolution2D(Get_filter(2), Kernel, name= "Dec_Block"+str(Block_Number)+"_Conv2", activation = 'relu', padding= "same", kernel_initializer= "he_normal")(Layer_1)
Layer_2 = BatchNormalization(name= "Dec_Block"+str(Block_Number)+"_Batch2")(Layer_2)

if Block_Number>2:
    # Defining Third Convolution layers with Batch Normalization for necessary Decoder Block
    Layer_3 = Convolution2D(Get_filter(3), Kernel, name= "Dec_Block"+str(Block_Number)+"_Conv3", activation = 'relu', padding= "same", kernel_initializer= "he_normal")(Layer_2)
    Layer_3 = BatchNormalization(name= "Dec_Block"+str(Block_Number)+"_Batch3")(Layer_3)
    return Layer_3

return Layer_2

# Building Encoder Block for Segnet with Various filters Sizes
Encoder_Conv1=Segnet_Encoder(1, 64, Inputs)
max_pool1,pool_indices_1 =tf.nn.max_pool_with_argmax(Encoder_Conv1, [1, 2, 2, 1], [1, 2, 2, 1], padding="VALID")

Encoder_Conv2=Segnet_Encoder(2, 128, max_pool1)
max_pool_2,pool_indices_2 =tf.nn.max_pool_with_argmax(Encoder_Conv2, [1, 2, 2, 1], [1, 2, 2, 1], padding="VALID")

Encoder_Conv3=Segnet_Encoder(3, 256, max_pool_2)
max_pool_3,pool_indices_3 =tf.nn.max_pool_with_argmax(Encoder_Conv3, [1, 2, 2, 1], [1, 2, 2, 1], padding="VALID")

Encoder_Conv4=Segnet_Encoder(4, 512, max_pool_3)
max_pool_4,pool_indices_4 =tf.nn.max_pool_with_argmax(Encoder_Conv4, [1, 2, 2, 1], [1, 2, 2, 1], padding="VALID")

Encoder_Conv5=Segnet_Encoder(5, 512, max_pool_4)
max_pool_5,pool_indices_5 =tf.nn.max_pool_with_argmax(Encoder_Conv5, [1, 2, 2, 1], [1, 2, 2, 1], padding="VALID")

# Building Decoder Block for Segnet with Various filters Sizes
Decoder_upsamp5 =MaxUnpooling2D(max_pool_5, pool_indices_5, batch_size, name="un_pool_5")
Decoder_Conv5=Segnet_Decoder(5, 512, Decoder_upsamp5)

Decoder_upsamp4 =MaxUnpooling2D(Decoder_Conv5, pool_indices_4, batch_size, name="un_pool_4")
Decoder_Conv4=Segnet_Decoder(4, 512, Decoder_upsamp4)

Decoder_upsamp3 =MaxUnpooling2D(Decoder_Conv4, pool_indices_3, batch_size, name="un_pool_3")
Decoder_Conv3=Segnet_Decoder(3, 256, Decoder_upsamp3)

Decoder_upsamp2 =MaxUnpooling2D(Decoder_Conv3, pool_indices_2, batch_size, name="un_pool_2")
Decoder_Conv2=Segnet_Decoder(2, 128, Decoder_upsamp2)

Decoder_upsamp1 =MaxUnpooling2D(Decoder_Conv2, pool_indices_1, batch_size, name="un_pool_1")
Decoder_Conv1=Segnet_Decoder(1, 64, Decoder_upsamp1)

# Final Convolution layer has number of classes as filter size followed by softmax Layer
Convolution_out = tf.keras.layers.Conv2D(n_labels, (3,3), name= "Final_Conv", activation = 'relu', kernel_initializer="he_normal", padding = 'same')(Decoder_Conv1)
Output=Activation('softmax', name="Softmax")(Convolution_out)

# Invoke to get Segnet model
model = Model(inputs=Inputs, outputs=Output, name="SEGNET")

return model

```

Training Segnet Model

In []:

```

# Get current Time
start_time = time()

# Defining Batch size and epoch
batch_size, epochs = 16, 30

# Defining tensorboard to store Training Information and filepath to store Segnet model
tensorboard, filepath = TensorBoard(log_dir=root+"logs/Segnet_{}".format(str(time())[5:10])), root+"Segne

```

```
t.best.hdf5"
```

```
# Defining steps_per_epoch and validation_steps for Training
steps_per_epoch, validation_steps = int((len(train_img_files1)+len(train_img_files2))/batch_size),int((len(val_img_files1)+len(val_img_files2))/batch_size)

# Invoke Segnet_Segmentation to get Segnet Model
Segnet_Model = Segnet_Segmentation((240,480,3), 7, batch_size,Kernel=(3,3), Pool_Filter=(2, 2), output_mode= "softmax")

# Compile Segnet Model
Segnet_Model.compile(optimizer = tf.keras.optimizers.Adam(0.0001), loss = 'categorical_crossentropy',metrics = ['accuracy',miou])

# Defining EarlyStopping with patience=5 and monitor='val_miou'
es = EarlyStopping(monitor='val_miou', mode='max', verbose=1, patience=5)

# Defining ModelCheckpoint with monitor as 'val_miou'
checkpoint = ModelCheckpoint(filepath, monitor='val_miou', verbose=2, save_best_only=True, mode='max')

# Defining ReduceLROnPlateau to reduce learning rate with patience=3
learning_rate_reduction = ReduceLROnPlateau(monitor='val_miou', patience=3, verbose=2, factor=0.2, min_lr=0.00001)

# Fit Segnet Model to start training
history=Segnet_Model.fit_generator(train_batch_generator(batch_size,epochs), steps_per_epoch=steps_per_epoch, epochs=epochs, verbose=1, validation_data=val_batch_generator(batch_size,epochs), validation_steps=validation_steps, callbacks=[learning_rate_reduction,checkpoint,es,tensorboard])

# Printing Time taken for Training
print("--- %s seconds ---" % (time() - start_time))
```

WARNING:tensorflow:From <ipython-input-7-862c63110dc2>:11: Model.fit_generator (from tensorflow.python.keras.engine.training) is deprecated and will be removed in a future version.

Instructions for updating:

Please use Model.fit, which supports generators.

Epoch 1/30

1/876 [.....] - ETA: 0s - loss: 2.1792 - accuracy: 0.1306 - miou: 0.0588WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/summary_ops_v2.py:1277: stop (from tensorflow.python.eager.profiler) is deprecated and will be removed after 2020-07-01.

Instructions for updating:

use `tf.profiler.experimental.stop` instead.

876/876 [=====] - ETA: 0s - loss: 0.6153 - accuracy: 0.7792 - miou: 0.4374

Epoch 00001: val_miou improved from -inf to 0.46177, saving model to /content/drive/My Drive/Segnet2.best.hdf5

876/876 [=====] - 2409s 3s/step - loss: 0.6153 - accuracy: 0.7792 - miou: 0.4374 - val_loss: 0.4914 - val_accuracy: 0.8156 - val_miou: 0.4618

Epoch 2/30

876/876 [=====] - ETA: 0s - loss: 0.3991 - accuracy: 0.8478 - miou: 0.5519

Epoch 00002: val_miou improved from 0.46177 to 0.54427, saving model to /content/drive/My Drive/Segnet2.best.hdf5

876/876 [=====] - 846s 966ms/step - loss: 0.3991 - accuracy: 0.8478 - miou: 0.5519 - val_loss: 0.4324 - val_accuracy: 0.8364 - val_miou: 0.5443

Epoch 3/30

876/876 [=====] - ETA: 0s - loss: 0.3383 - accuracy: 0.8706 - miou: 0.6196

Epoch 00003: val_miou improved from 0.54427 to 0.58790, saving model to /content/drive/My Drive/Segnet2.best.hdf5

876/876 [=====] - 848s 968ms/step - loss: 0.3383 - accuracy: 0.8706 - miou: 0.6196 - val_loss: 0.3676 - val_accuracy: 0.8575 - val_miou: 0.5879

Epoch 4/30

876/876 [=====] - ETA: 0s - loss: 0.3009 - accuracy: 0.8847 - miou: 0.6548

Epoch 00004: ReduceLROnPlateau reducing learning rate to 1.9999999494757503e-05.

Epoch 00004: val_miou did not improve from 0.58790

876/876 [=====] - 837s 956ms/step - loss: 0.3009 - accuracy: 0.8847 - miou: 0.6548 - val_loss: 0.3769 - val_accuracy: 0.8627 - val_miou: 0.5842

Epoch 5/30

876/876 [=====] - ETA: 0s - loss: 0.2540 - accuracy: 0.9030 - miou: 0.6982

Epoch 00005: val_miou improved from 0.58790 to 0.63156, saving model to /content/drive/My Drive/Segnet2.best.hdf5

876/876 [=====] - 844s 964ms/step - loss: 0.2540 - accuracy: 0.9030 - miou: 0.6982 - val_loss: 0.3161 - val_accuracy: 0.8804 - val_miou: 0.6316

Epoch 6/30

876/876 [=====] - ETA: 0s - loss: 0.2289 - accuracy: 0.9131 - miou: 0.7216

Epoch 00006: val_miou improved from 0.63156 to 0.63321, saving model to /content/drive/My Drive/Segnet2.best.hdf5

876/876 [=====] - 845s 964ms/step - loss: 0.2289 - accuracy: 0.9131 - miou: 0.7216 - val_loss: 0.3325 - val_accuracy: 0.8792 - val_miou: 0.6332

Epoch 7/30

876/876 [=====] - ETA: 0s - loss: 0.2110 - accuracy: 0.9204 - miou: 0.7414

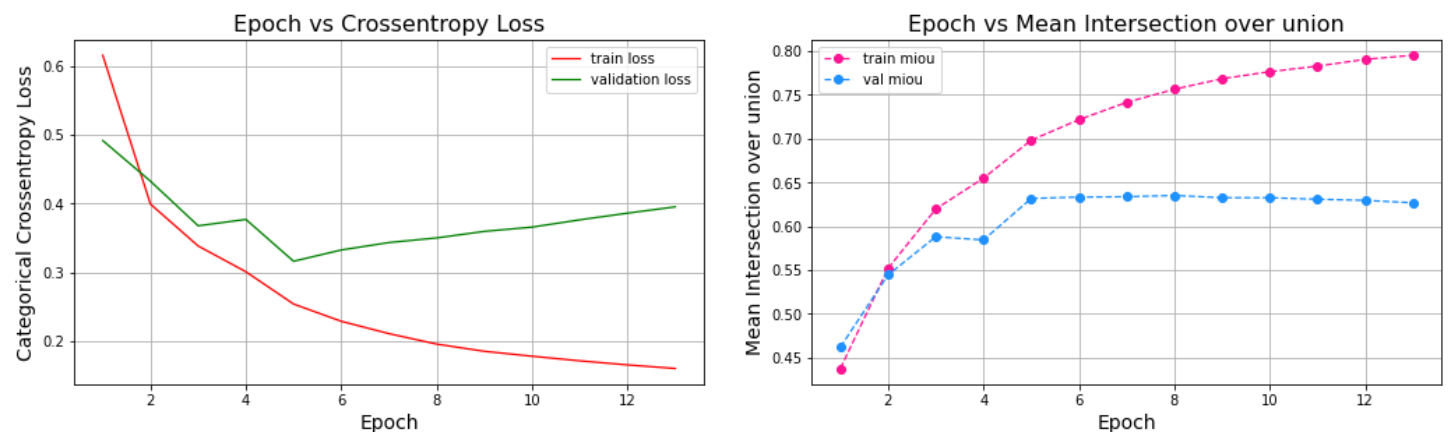
Epoch 00007: ReduceLROnPlateau reducing learning rate to 1e-05.

```
Epoch 00007: val_miou improved from 0.63321 to 0.63355, saving model to /content/drive/My Drive/Segnet2.be
st.hdf5
876/876 [=====] - 843s 962ms/step - loss: 0.2110 - accuracy: 0.9204 - miou: 0.741
4 - val_loss: 0.3433 - val_accuracy: 0.8792 - val_miou: 0.6336
Epoch 8/30
876/876 [=====] - ETA: 0s - loss: 0.1958 - accuracy: 0.9265 - miou: 0.7563
Epoch 00008: val_miou improved from 0.63355 to 0.63497, saving model to /content/drive/My Drive/Segnet2.be
st.hdf5
876/876 [=====] - 839s 958ms/step - loss: 0.1958 - accuracy: 0.9265 - miou: 0.756
3 - val_loss: 0.3500 - val_accuracy: 0.8783 - val_miou: 0.6350
Epoch 9/30
876/876 [=====] - ETA: 0s - loss: 0.1853 - accuracy: 0.9306 - miou: 0.7684
Epoch 00009: val_miou did not improve from 0.63497
876/876 [=====] - 839s 957ms/step - loss: 0.1853 - accuracy: 0.9306 - miou: 0.768
4 - val_loss: 0.3595 - val_accuracy: 0.8807 - val_miou: 0.6325
Epoch 10/30
876/876 [=====] - ETA: 0s - loss: 0.1782 - accuracy: 0.9334 - miou: 0.7764
Epoch 00010: val_miou did not improve from 0.63497
876/876 [=====] - 834s 952ms/step - loss: 0.1782 - accuracy: 0.9334 - miou: 0.776
4 - val_loss: 0.3657 - val_accuracy: 0.8793 - val_miou: 0.6324
Epoch 11/30
876/876 [=====] - ETA: 0s - loss: 0.1715 - accuracy: 0.9358 - miou: 0.7828
Epoch 00011: val_miou did not improve from 0.63497
876/876 [=====] - 838s 956ms/step - loss: 0.1715 - accuracy: 0.9358 - miou: 0.782
8 - val_loss: 0.3763 - val_accuracy: 0.8798 - val_miou: 0.6307
Epoch 12/30
876/876 [=====] - ETA: 0s - loss: 0.1656 - accuracy: 0.9381 - miou: 0.7905
Epoch 00012: val_miou did not improve from 0.63497
876/876 [=====] - 835s 953ms/step - loss: 0.1656 - accuracy: 0.9381 - miou: 0.790
5 - val_loss: 0.3860 - val_accuracy: 0.8795 - val_miou: 0.6294
Epoch 13/30
876/876 [=====] - ETA: 0s - loss: 0.1604 - accuracy: 0.9401 - miou: 0.7952
Epoch 00013: val_miou did not improve from 0.63497
876/876 [=====] - 838s 957ms/step - loss: 0.1604 - accuracy: 0.9401 - miou: 0.795
2 - val_loss: 0.3951 - val_accuracy: 0.8783 - val_miou: 0.6267
Epoch 00013: early stopping
--- 12524.988194227219 seconds ---
```

Segnet Training Results

In []:

```
# training_result
plot_training_result(history)
```



- The Lowest value of Validation Categorical Crosseentropy is 0.3161 which is at epoch-5 as above in the Graph.
- The Best Value of Validation Mean Intersection Over Union is 0.6349 which is at epoch-8 as above in the Graph.
- Keras callback ModelCheckpoint is used to save the best Model during Training to avoid overfitting.

Segnet Prediction on Train Data

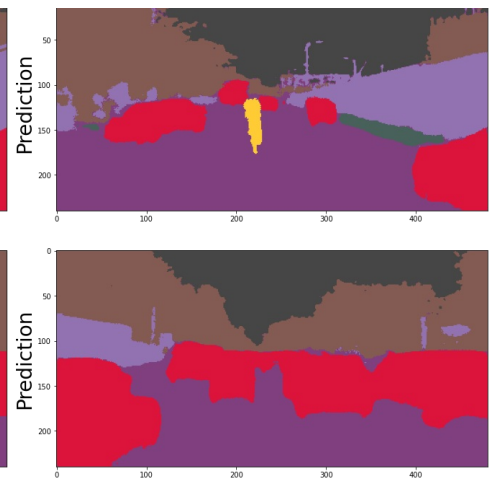

In []:

```
# Train Prediction
Miou, Accuracy, cf_matrix = predict_for("Train_data")
```

Total number of samples in Train_data : 10016

Few Segmentation Samples:>>>





Printing Results:>>

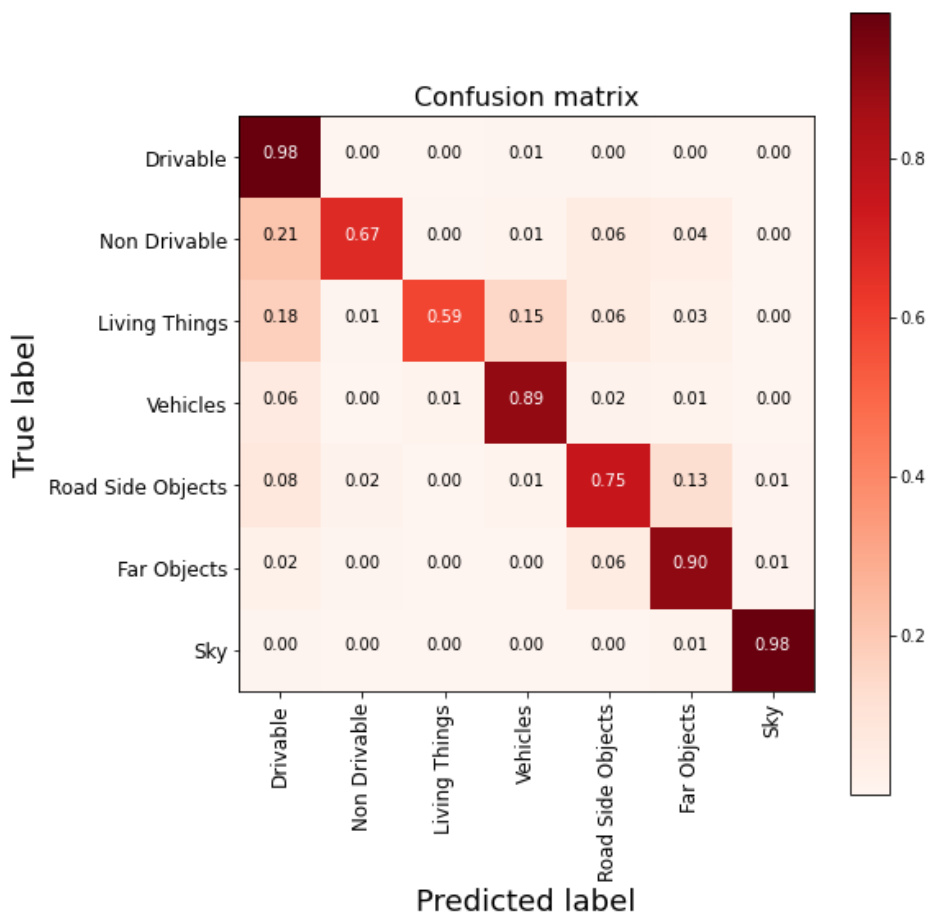
MIOU Score

MIOU Score: 0.6258

Accuracy Score

Accuracy Score: 0.9154

Confusion Matrix



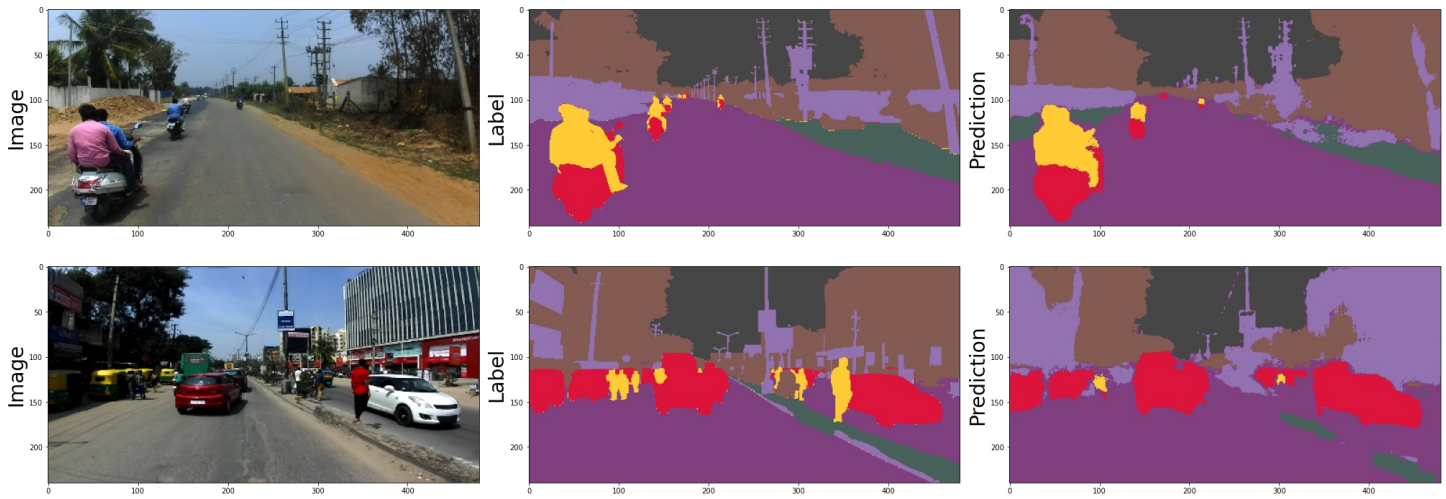
Segnet Prediction on Validation Data

In []:

```
# Validation Prediction
Miou, Accuracy, cf_matrix = predict_for("Val_data")
```

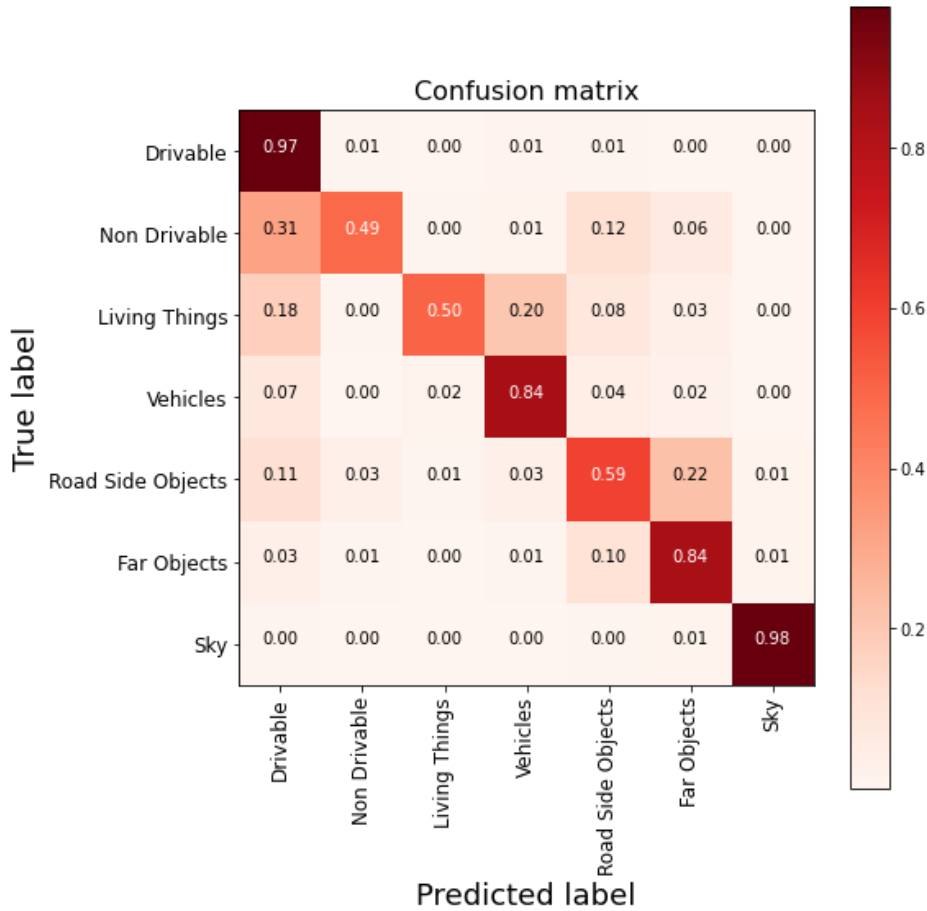
Total number of samples in Val_data : 2036

Few Segmentation Samples:>>>



Printing Results:>>

```
-----
|      MIOU Score      |
|-----|
|
| MIOU Score: 0.5643
|
|-----|
|    Accuracy Score    |
|-----|
|
| Accuracy Score: 0.8698
|
|-----|
|   Confusion Matrix   |
|-----|
```



Segnet Prediction on Test Data


```
In [ ]:
# Test Prediction
```



```
Miou, Accuracy, cf_matrix = predict_for("Test_data")
```

Total number of samples in Test_data : 4011

Few Segmentation Samples:>>>



Printing Results:>>

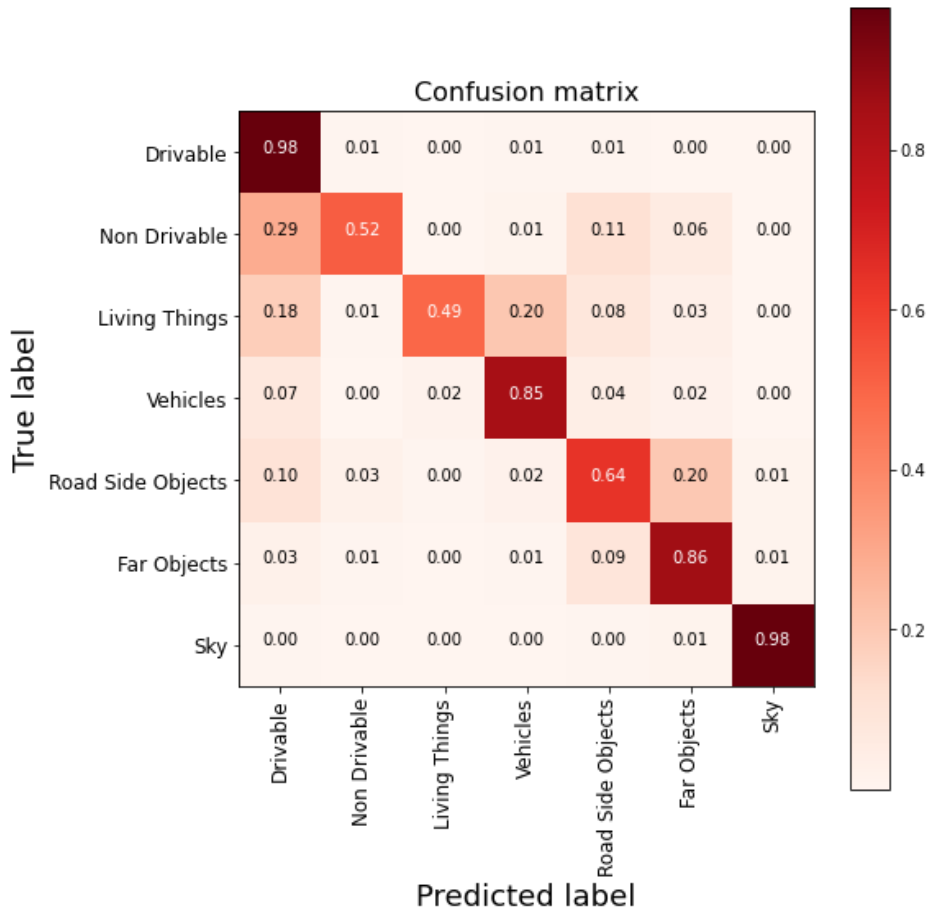
```
-----  
|      MIOU Score      |  
-----
```

MIOU Score: 0.5747

```
-----  
|   Accuracy Score   |  
-----
```

Accuracy Score: 0.8814

```
-----  
|  Confusion Matrix  |  
-----
```



Pretty Tabel

```
In [ ]:  
  
# https://ptable.readthedocs.io/en/latest/tutorial.html  
print("\n\t\t\tPerformance Table")  
from prettytable import PrettyTable  
T = PrettyTable()  
T.field_names = ["SegNet", "MIOU", "Accuracy"]  
T.add_row(["Train ", "0.6258", "0.9154"])  
T.add_row(["-----", "-----", "-----"])
```

```
T.add_row(["Validation ", "0.5643", "0.8698"])
T.add_row([" ----- ", "-----", "-----"])
T.add_row(["Test ", " 0.5747", "0.8814"])
print(T)
```

Performance Table			
-----+	-----+	-----+	-----+
SegNet	MIOU	Accuracy	
-----+	-----+	-----+	-----+
Train	0.6258	0.9154	
-----+	-----+	-----+	-----+
Validation	0.5643	0.8698	
-----+	-----+	-----+	-----+
Test	0.5747	0.8814	
-----+	-----+	-----+	-----+

****Conclusion:****

- SegNet is an Encoder-Decoder Network which is efficient both in terms of memory and computational time
- SegNet focuses on the need to map low-resolution features to input resolution for Image Segmentation.
- SegNet stores the max-pooling indices of the feature maps and uses them in its decoder network to achieve good performance as above.
- The SegNet architecture achieves relatively good performance on Image segmentation when compared to other segmentation models.
- More performance can be obtained by training Models with data in high resolution with more powerful hardware resources